

# PERFORMANCE MANAGEMENT

## Practical 2

# Introduction to Performance Management:

## What is Performance Management ?

Performance management in the context of a database refers to the practice of monitoring, optimizing, and maintaining the efficiency and responsiveness of a database system.

### Impacts

It is crucial for organizations relying on databases because poor performance can lead to slow query responses, application downtime, and ultimately impact business operations and customer satisfaction.

# Performance Metrics

Performance metrics are measurements used to evaluate the effectiveness and efficiency of a database system.

## Common metrics include:

- **Response Time:** The time it takes for a database operation (e.g., a query) to complete.
- **Throughput:** The number of transactions or operations processed per unit of time.
- **Resource Utilization:** Monitoring CPU, memory, disk, and network usage.
- **Concurrency:** Assessing the level of concurrent users or transactions.



# Identifying Performance Bottlenecks

Performance bottlenecks are points in the database system where resources are constrained, causing a slowdown in performance.

- **CPU Bottleneck:** A CPU bottleneck in database performance management occurs when the central processing unit (CPU) of the database server becomes a limiting factor in the overall system's ability to process and execute database operations efficiently.

Several factors can contribute to CPU bottlenecks in database performance

# Identifying Performance Bottlenecks (Cont.)

## CPU Bottleneck

**Query Complexity:** Complex and resource-intensive database queries, such as those involving multiple joins, subqueries, or extensive sorting and filtering, can put a significant load on the CPU. Poorly optimized or inefficient queries are more likely to cause CPU bottlenecks.

**Insufficient Indexing:** Inadequate or missing indexes can lead to full table scans and increased CPU utilization. Proper indexing helps the database engine retrieve data more efficiently, reducing CPU overhead.

**Data Volume:** As the amount of data in the database grows, the CPU may struggle to process large datasets, especially when performing aggregations or calculations on extensive data.

**Concurrency:** High levels of concurrent user activity can lead to CPU contention, as multiple queries compete for processing resources. This is common in systems with many simultaneous users or transactions.

**Inefficient Code:** Poorly written stored procedures, triggers, or application code can consume more



# Identifying Performance Bottlenecks (Cont.)

## CPU Bottleneck

**Hardware Limitations:** The CPU's processing power may become a bottleneck if the hardware is not adequately scaled to meet the database's demands. Adding more CPU cores or upgrading to faster CPUs may be necessary in such cases.

**Lack of Parallelism:** Some database systems can benefit from parallel processing, where multiple CPU cores work on different parts of a query simultaneously. If the database management system (DBMS) or queries are not configured to take advantage of parallelism, it can result in CPU bottlenecks.

**Locking and Blocking:** Contentious locking and blocking issues can cause queries to wait for resources, leading to increased CPU usage as they remain active while waiting for access to data.

**Resource-Intensive Applications:** Other applications running on the same server as the database may consume CPU resources, leaving fewer available for the DBMS.

**Inadequate Hardware Resources:** If the database server lacks sufficient RAM, storage I/O capacity, or network bandwidth, it can indirectly cause CPU bottlenecks by forcing the CPU to work harder to compensate for these deficiencies.

**Background Processes:** Database systems often have background processes, such as backups, maintenance, or replication, which can consume CPU resources. Poorly scheduled or resource-intensive background processes

# Identifying Performance Bottlenecks (Cont.)

- **Memory Bottleneck:** When there's insufficient memory for caching frequently accessed data.
- **I/O Bottleneck:** When the storage subsystem (disk) cannot keep up with read/write requests.
- **Lock Contention:** When multiple sessions are waiting for locks, leading to delays.

**Example:** If a database query is slow due to excessive disk I/O, optimizing I/O operations can alleviate the bottleneck.



# SQL Tuning

SQL tuning involves optimizing SQL queries to improve their performance.

## Techniques

- **Index Usage:** Adding appropriate indexes to tables to speed up data retrieval.
- **Query Rewrite:** Rewriting queries to use more efficient SQL constructs.
- **SQL Profiling:** Analyzing query execution plans and making adjustments.

**Example:** Rewriting a complex JOIN query to use subqueries or optimizing an index to speed up SELECT queries.



# Database Tuning Approaches

Two main approaches to database tuning

- **Reactive Tuning**
- **Proactive Tuning**

# Database Tuning Approaches

## Reactive Tuning

Reactive tuning is a responsive approach to database performance management. It involves identifying and addressing performance issues as they occur or are reported by users.

- **Characteristics:**

- **Issue Identification:** In reactive tuning, problems are often discovered when users experience slow query responses, system errors, or downtime.
- **Immediate Action:** Database administrators respond to issues on an ad-hoc basis, implementing fixes or workarounds to resolve the problem at hand.
- **Short-Term Focus:** The primary goal is to resolve the current performance problem to restore normal database operations.



# Database Tuning Approaches

## Reactive Tuning

### Examples:

- If a critical application suddenly experiences slow query performance, administrators might identify and optimize the poorly performing SQL queries.
- When a database server crashes due to resource exhaustion, administrators may quickly restart the server and allocate additional resources.

# Database Tuning Approaches

## Proactive Tuning

Proactive tuning is a preventative approach to database performance management. It involves taking measures to optimize the database in advance, minimizing the likelihood of performance problems.

### Characteristics:

- **Continuous Monitoring**: Proactive tuning relies on ongoing monitoring of database performance, even when there are no apparent issues.
- **Performance Baselines**: Database administrators establish performance baselines to understand normal operation and detect deviations.
- **Optimization Strategies**: Administrators employ strategies such as indexing, query optimization, and resource allocation based on anticipated growth and usage patterns.
- **Capacity Planning**: Capacity planning helps ensure that the database has the necessary resources to handle future demands.



# Database Tuning Approaches

## Proactive Tuning

### Examples

- Regularly analyzing query execution plans and optimizing SQL statements to prevent performance bottlenecks.
- Conducting load testing to simulate heavy user traffic and identifying potential scalability issues in advance.
- Allocating additional memory or storage resources before they become critical, based on historical usage patterns.

# Memory Management

Memory management in Oracle involves allocating and managing memory structures for database operations.

## Key memory components:

- **Buffer Cache:** Caches frequently accessed data blocks in memory to reduce disk I/O.
- **Shared Pool:** Stores SQL statements, parsed execution plans, and other shared resources.
- **Program Global Area (PGA):** Memory for sorting, hash joins, and other user session-specific data.

**Example:** Configuring an appropriate buffer cache size to minimize disk I/O for read-heavy applications.



# I/O Optimization

- I/O optimization focuses on minimizing input/output operations to storage devices.

## Strategies

- **Tablespace and Data File Placement:** Organizing data files to reduce seek time.
- **RAID Configurations:** Implementing Redundant Array of Independent Disks for fault tolerance and performance.
- **Storage Optimization:** Using solid-state drives (SSDs) for high-speed I/O.

**Example:** Implementing RAID 10 for a high-transaction database to improve both performance and fault tolerance.

# Concurrency Management

Concurrency management deals with ensuring multiple users can access and modify the database concurrently without conflicts.

## Techniques

- **Locking:** Managing locks to prevent data inconsistencies.
- **Latches:** Controlling access to in-memory data structures to avoid contention.
- **Isolation Levels:** Defining the level of isolation between transactions.

**Example:** Implementing row-level locking to allow multiple users to update different rows simultaneously.



# Automatic Performance Tuning

Oracle provides automatic tuning features that assist in identifying and resolving performance issues.

## Examples:

- **Automatic SQL Tuning:** Automatically optimizes SQL statements for better performance.
- **Automatic Memory Management:** Adjusts memory allocations based on system needs.

**Example:** The Automatic SQL Tuning Advisor can recommend index creation or query rewrites to improve SQL performance without manual intervention.

# Performance Testing and Benchmarking

Performance testing involves evaluating the behavior and capabilities of a database system under various conditions to ensure it meets performance requirements.

## Types of performance testing

- **Load Testing:** Evaluates how the system performs under expected load levels.
- **Stress Testing:** Tests the system's performance at or beyond its limits to identify breaking points.
- **Benchmarking:** Compares the system's performance against industry standards or competitors.

**Example:** Load testing a web application backed by an Oracle database to determine how many concurrent users it can handle without degrading performance.



# Capacity Planning

Capacity planning involves forecasting future resource requirements to ensure the database can handle increased data volumes and user loads.

## Key considerations

- **Data Growth:** Estimate how quickly data will grow over time.
- **Workload Growth:** Predict increases in user activity and transactions.
- **Resource Scalability:** Plan for scaling CPU, memory, storage, and network resources.

**Example:** A retail company estimates that its database will need to accommodate a 20% increase in transaction volume during the holiday season and plans to scale its server resources accordingly.